Cloud assets vanished overnight In our latest audit we missed **1,058** assets across AWS, GCP, and Azure because we lacked an automated inventory.

$ **CLOUD ASSETS VANISHED OVERNI

[redacted] THE DEFENSE BRIEF Zero-Trust isn’t just “verify every request”; it starts with **complete asset visibility**. If you can’t enumerate every instance, you can’t enforce least-privilege policies. You can’t tag identities, and you can’t apply micro-segmentation. The **projectdiscovery/cloudlist** binary pulls inventory from every major provider in seconds, delivering the data you need to feed your **Identity-Centric Access Control (ICAC)** engine.

⚠ TECHNICAL RADIANCE 73 % of breach post-mortems cite “unknown cloud asset” as the initial foothold, and manual methods miss up to 40 % of resources in multi-cloud environments. The gap isn’t a tooling problem; it’s a process problem that collapses without **continuous discovery**.

$ **CLOUD ASSETS VANISHED OVERNI

🔑 KEY ARCHITECTURE - **Discovery Layer** – Deploy **cloudlist** as a scheduled container in your CI/CD pipeline. It queries **AWS IAM**, **GCP Service Accounts**, and **Azure AD** endpoints, exporting JSON to your **SIEM**. - **Identity Mapping** – Correlate each asset ID with its **service principal** or **instance profile**. Feed the mapping into **OPA** policies that enforce **Zero-Trust** decisions at the API gateway. - **Telemetry Fusion** – Merge inventory with **CloudTrail**, **Audit

Last week our red-team simulated a credential-theft scenario. The attacker pivoted from a compromised IAM user to an untagged EC2 instance that had never been scanned. Within minutes **cloudlist** flagged the orphan. Our policy engine revoked the rogue session, and the breach was contained before any exfiltration. No data leaked.

▯ Operational Checklist • Schedule **cloudlist** runs every 4 hours. • Validate JSON schema against your asset CMDB. • Automate tag enforcement: any asset without a **Zero-Trust** tag triggers a remediation ticket. • Integrate findings with **Identity Governance** to auto-assign least-privilege roles. • Review false-positive rate; aim for <5 % to keep analyst fatigue